



MONASH
University

MONASH
INFORMATION
TECHNOLOGY

INTERPROCESS COMMUNICATION

LECTURE 10 / FIT2100 / SEMESTER 2
2019

WEEK 11



INTERPROCESS COMMUNICATION

INTRODUCTION

LEARNING OUTCOMES

- Discuss various mechanisms for interprocess communication (IPC).
- Understand the implications of different IPC mechanisms for **synchronization** and **concurrency**.

READING

- Stallings:
 - Chapter 5 sections 5.6, 5.7.
 - Chapter 6 sections 6.7, 6.8.
- Further reading: Curry, *Unix Systems Programming for SVR4*
 - Chapter 13 – IPC, Chapter 10 – Signals

WHAT IS IPC?

INTERPROCESS COMMUNICATION

COMMUNICATION BETWEEN TWO OR MORE PROCESSES

- Multiple applications/utilities often need to communicate
- Communication may happen over a network, but often communication happens among processes running on the same computer.

APPROACHES TO BE DISCUSSED

- Signals
- Shared memory
- Streams (pipes)
- Sockets (client/server)
- Message queues.

WHY DO WE NEED IPC?

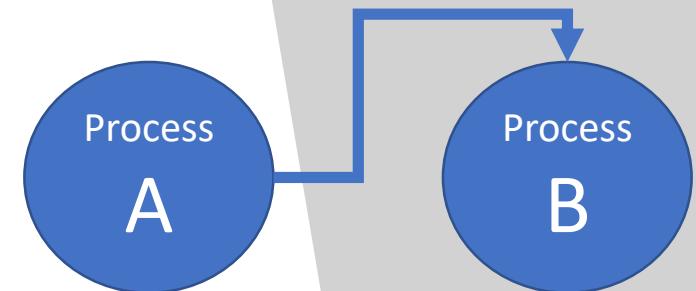
LET'S COMPARE PROCESSES TO THREADS...

THREADS

- Memory is ‘shared’ between threads
 - Well actually, there is only one memory space
 - The threads are all part of the same process
 - This is only communication **within** a process

PROCESSES

- Often, multiple related or unrelated processes need to exchange data
 - Most common case: a program often streams output to a console utility using a ‘**print**’ statement.
- Different processes do not typically share the same memory space.
- The OS provides a number of communication mechanisms here.



SIGNALS

THE MOST PRIMITIVE KIND OF IPC



Image credit: Original mikz, Creative Commons:
BY-SA 4.0

UNIX SIGNALS

SIGNALS ARE LIKE INTERRUPTS FOR USER PROCESSES

USED TO INFORM A PROCESS OF AN ASYNCHRONOUS EVENT

- A signal is ‘delivered’ by updating a field (bit flag) in the process table for the process that receives the signal.
- There are no priorities for signals: all treated equally.
- A signal has a signal number and that’s it. This number identifies the **type** of signal sent.

A PROCESS MAY RESPOND TO A SPECIFIED SIGNAL NUMBER BY...

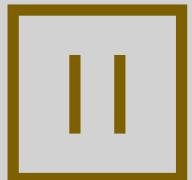
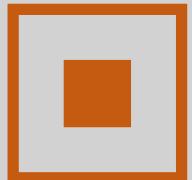
- Jumping into a signal-handler function, or
- Choosing to ignore a signal, or,
- Performing the operating system’s default action for that signal (e.g. process termination).

SOME COMMON SIGNALS

FOR EACH SIGNAL NUMBER, A CONSTANT IS DEFINED IN `<signal.h>`

SIGNALS USED FOR PROCESS MANAGEMENT

- **SIGINT**: The *interrupt* signal that is sent when you press `Ctrl+C` in a terminal window.
 - **Default action**: terminate the process.
- **SIGTSTP**: The *stop* signal that is sent when you press `Ctrl+Z` in a terminal window
 - **Default action**: suspend (pause) the process
 - In `bash`, the `fg` command can be used to resume a stopped process.



SOME SIGNALS CANNOT BE HANDLED OR IGNORED BY PROCESSES.

- **SIGKILL**: Used to *kill* a process without any way for the process to ignore the signal or handle it in a different way
 - **Default (only) action**: terminate the process.

For more signals, check the man page: `man 7 signal`

DEALING WITH SIGNALS

UNIX SYSTEM CALLS (defined in <signal.h>)

HANDLING A SIGNAL

- The handler must be set up **before** the signal arrives:

```
sighandler_t signal(int signal_number, sighandler_t handler)
```

- The value of **handler** may be a pointer to a handler function to be called, or **SIG_IGN** (ignore signal) or **SIG_DFL** (reset to default action).
- Refer to man page: `man 2 signal` for more information.

SENDING A SIGNAL TO ANOTHER PROCESS

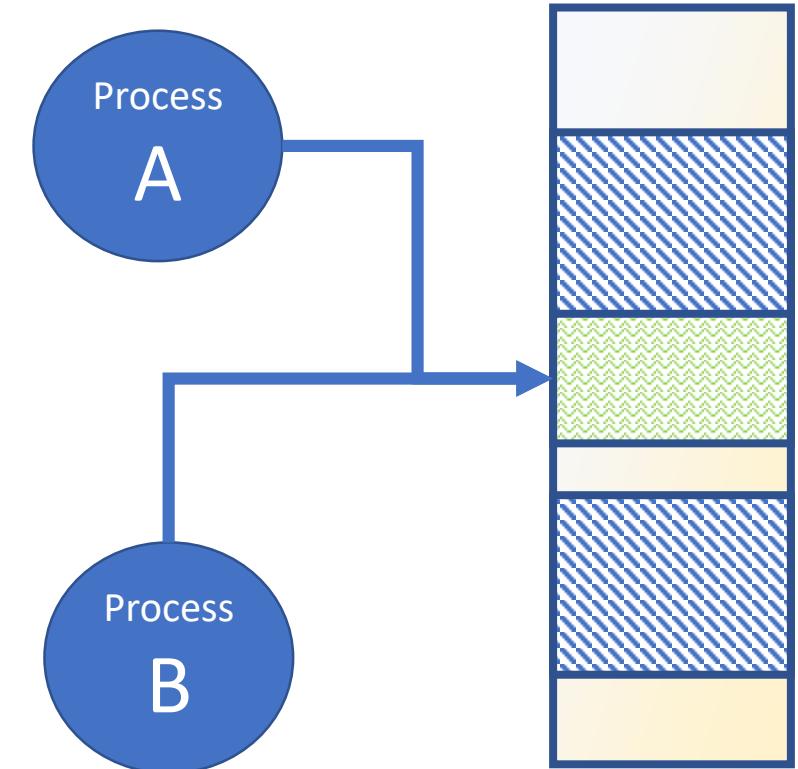
- You only need the PID and the signal type you wish to send:

```
int kill(pid_t process, int signal_number)
```

- Why such a morbid name, ‘kill’?
 - Originally the signalling system was only used to terminate processes.
 - Nowadays, a variety of harmless signals can be sent too!

SHARED MEMORY (SHM)

– AN IPC APPROACH THAT WORKS THROUGH THE VIRTUAL MEMORY SYSTEM



SHARED MEMORY

FROM LAST WEEK'S LECTURE...

A SEGMENT IN VITUAL MEMORY THAT IS ALLOCATED TO MULTIPLE PROCESSES

- A memory segment may be shared among multiple processes
- Like with multi-threading, multiple processes can access each other's memory
- Unlike with multi-threading, it does not happen automatically.
 - Must manually request a new block of shared memory to be allocated.
 - Required number of bytes must be specified.
 - Must manually assign a data structure (e.g. **struct** or **array**) to the shared block of bytes.
- There is no concurrency protection offered. Mutual exclusion mechanisms like semaphores must be used.
- Using multiple processes with shared memory can be **more robust** than using multiple threads. e.g. if one process crashes, the parent process can restart it.
 - e.g. tabs in Firefox web browser and Google Chrome

TWO WAYS TO USE SHARED MEMORY

LINUX SUPPORTS TWO TRADITIONAL UNIX APPROACHES

SYSTEM V STYLE (FROM OLD SVR4 UNIX SYSTEMS)

- Library functions in `<sys/shm.h>` include `shmget(...)` and `shmat(...)`
- Processes are required to have knowledge of a common **key** value in order to get access to the same segment.
- The key value is an integer that might be stored in a file for other processes to look up, or simply generated prior to forking a child process.

POSIX STYLE (NEWER ‘STANDARD’ APPROACH)

- Library functions in `<sys/mman.h>` include `shm_open(...)` and `mmap(...)`
- A **filename** in the filesystem is associated with a shared memory segment.
- When opened, the **file descriptor** can be read or written to just like a file.
- Can also be mapped to a logical memory address using `mmap(...)`

SHARED MEMORY APPROACHES

SYSTEM V VS POSIX



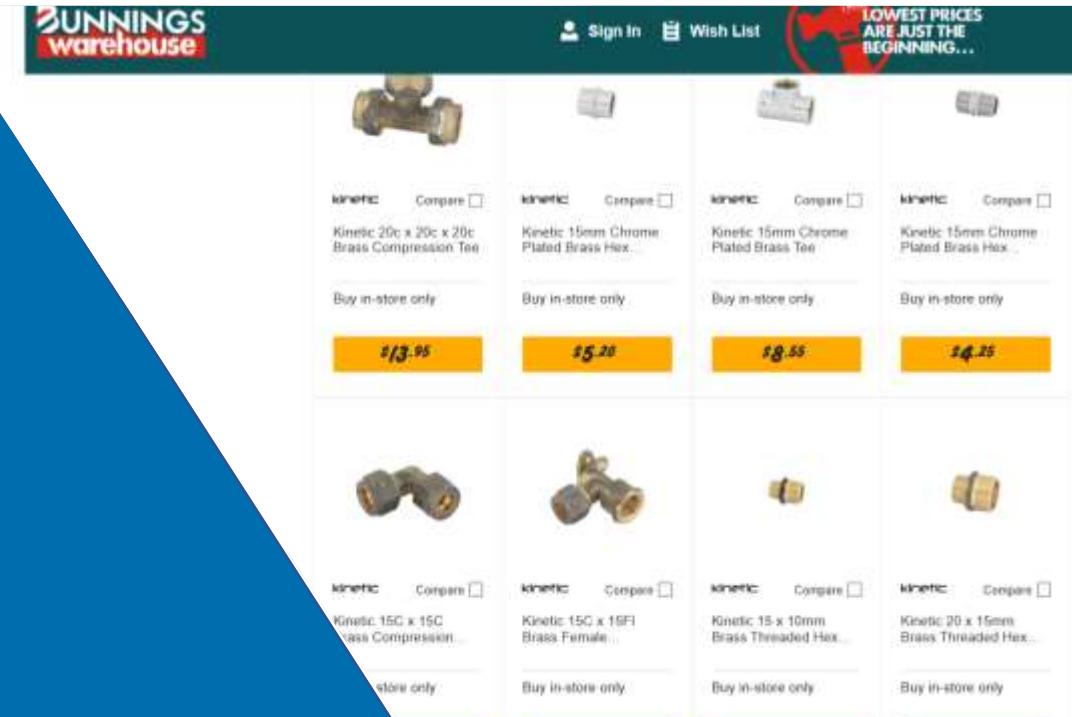
BEWARE!

- For either shared memory approach, there is **no** built-in concurrency protection
 - Shared memory has the same hazards as multithreading.
- Mutual exclusion mechanisms such as **semaphores**, **mutexes**, etc. must be used when accessing shared data.
- These mutual exclusion resources can be stored in the shared memory itself.

STREAM-BASED IPC

UNNAMED PIPES AND FIFOs

– AN APPROACH TO IPC THAT WORKS THROUGH THE KERNEL DIRECTLY



WHAT ARE PIPES?

BYTES GO IN ONE END, COME OUT AT THE OTHER END

A BUFFER MAINTAINED WITHIN THE KERNEL

- A pipe is a resource where information written in at one ‘end’ can be read out at the other ‘end’.
- Slower than SHM since all data must pass through the kernel.
- A pipe is a bit like a file but with **two** file descriptors. One file descriptor is for reading only, while the other is for writing only.

A PIPE ENFORCES SYNCHRONISATION BETWEEN READER AND WRITER

- As long as only one process reads, and only one process writes to pipe, concurrency protection is guaranteed.
 - First-in-first-out behavior

The reader cannot read past what has been written, and the writer cannot overwrite data in the pipe before it has been read

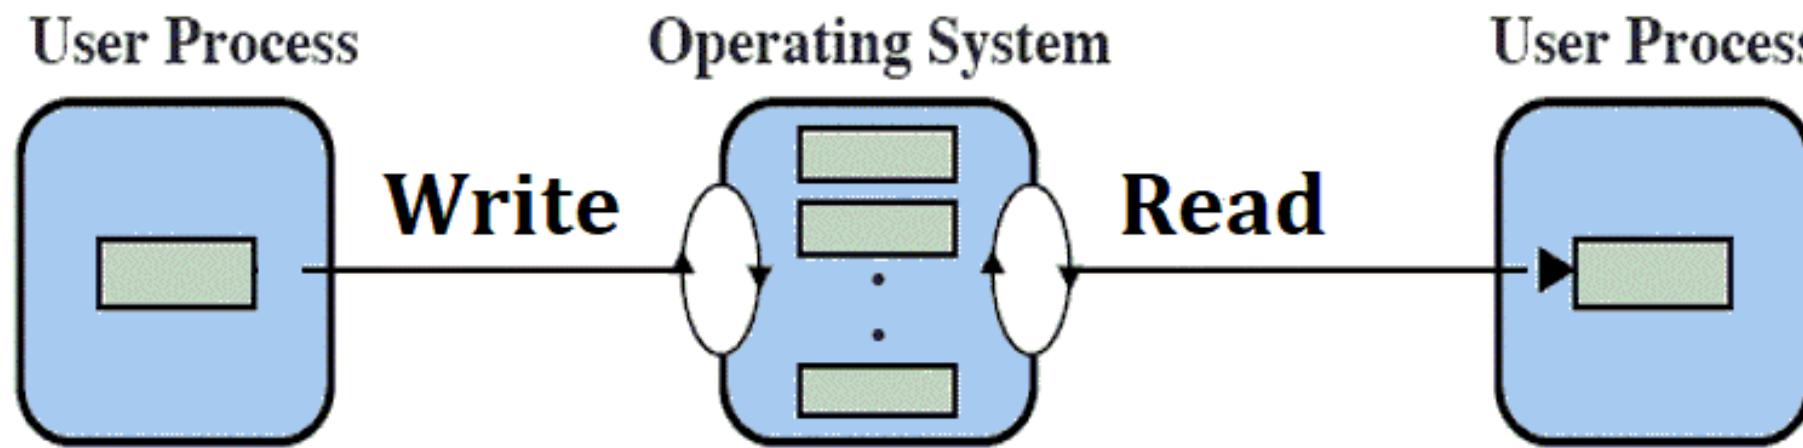
- The kernel implements a **circular buffer** to make this happen.

CIRCULAR BUFFERS (QUEUES)

RECALL: LECTURE 3 (WEEK 3)

A CIRCULAR BUFFER IS A SET OF BUFFERS ACCESSED IN A ‘CIRCULAR’ WAY

- The stream of bytes spans multiple buffers (character arrays)
- At any one time, one buffer is selected for reading while another is being written to.
- **Writer** fills buffers while the **reader** empties them, before moving on to the next buffer in the circle.
- If the reader ‘catches up’ to the writer, further reads will **block** until the next buffer becomes available (has been written to).
- If the writer **catches up** to the reader, further writes will **block** until the next buffer becomes available for writing (has been completely read).



TWO KINDS OF PIPES

UNNAMED AND NAMED

UNNAMED PIPES

- Can only be shared between related processes.
- Created using the `pipe` system call.
 - Creates a `pipe` and allocates **two file descriptors**.
 - After a `fork`, the child will inherit the parent's open file descriptors.
 - So one process can write to the pipe to talk to the other process.

NAMED PIPES (also known as FIFOs)

- A special kind of pipe with a **filename** in the filesystem.
- Can be used between unrelated processes.
- A special 'file' created using the `mknod` system call.
- One process then opens the 'file' read-only, while the other process opens the same file 'write-only'.

SOCKETS

– ONE SERVER PROCESS
ACCEPTS CONNECTIONS FROM
MULTIPLE CLIENT PROCESSES



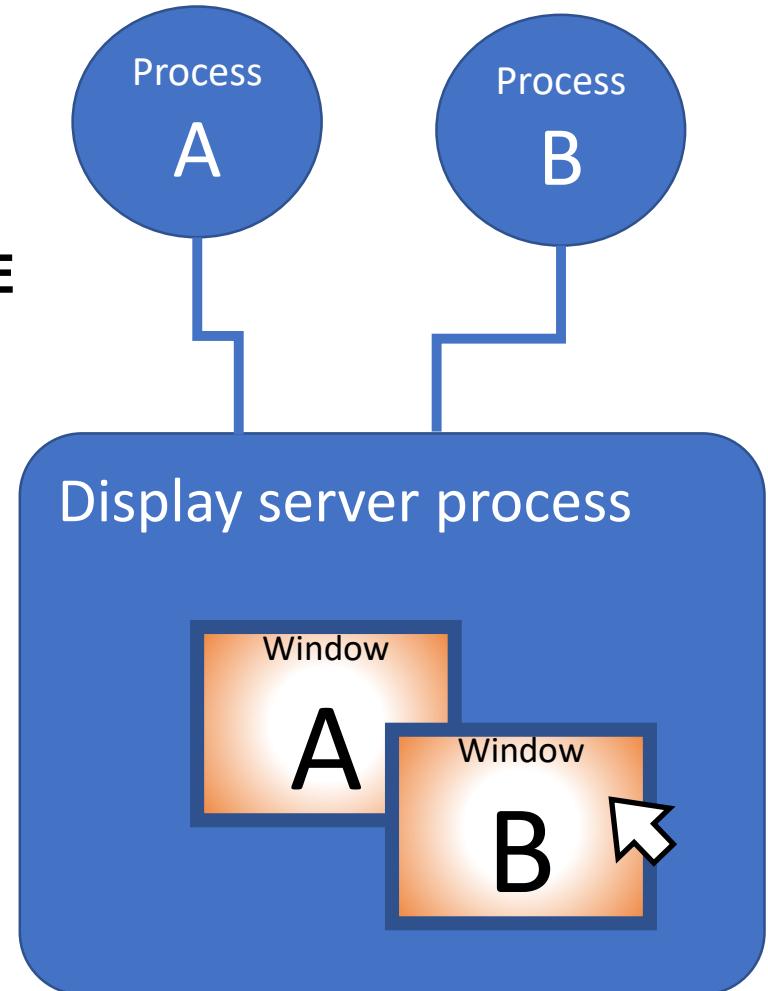
EXAMPLE

DISPLAY SERVER UTILITY

ONLY ONE PROCESS CAN CONTROL THE DISPLAY HARDWARE AT A TIME

How does a desktop system like **Windows** work?

- One utility called the **display server** takes responsibility for drawing all graphics on the screen.
- Applications (**clients**) connect to the display server and send instructions for what should be drawn on the screen. (Through an API.)
- The display server interprets instructions from different applications to produce a montage of **windows** on the screen.
- When the user interacts with the mouse and keyboard, the display server figures out which **client** to send the information back to.



If all processes were allowed to write to the display directly, the result would be a mess!

SOCKETS IN LINUX

'UNIX DOMAIN SOCKETS'

HOW IS A SOCKET CONNECTION CREATED?

1. The server process creates a **socket**, which is mapped to a **filename** in the filesystem.
2. Client processes must **connect** to this file to reach the server.
3. The server receives information about a new connection request, and chooses to **accept** the connection.
4. When accepted, the OS provides an open file descriptor resource to both the client and server process.
 - Just like a pipe, a socket connection enforces FIFO behaviour when it is being read and written to concurrently.
 - Unlike a pipe, a socket connection can send data in both directions.
 - The kernel provides **two** circular buffers. One for bytes sent from the client to the server, and another for bytes sent from the server to the client.

MESSAGE QUEUES

-A SYSTEM FOR SENDING FIXED-
LENGTH MESSAGES BETWEEN
PROCESSES



WHAT IS A MESSAGE QUEUE?

MORE STRUCTURED MESSAGE PASSING

A RESOURCE PROVIDED BY THE KERNEL

Once a queue has been created, multiple processes can send messages into the queue, or read them out of the queue.

Multiple processes can share a single queue.

MESSAGES COME WITH ATTACHED INFORMATION

By defining a set of different **message types** (assigning a different integer to different kinds of messages), a process can request the type of message it is waiting for to be delivered next.

Unlike streams (pipes, etc.) a message queue always serves a **whole message** at a time. It does not work on a per-character basis.

MESSAGE FORMAT

MESSAGES ARE SENT AS FIXED-LENGTH CHUNKS

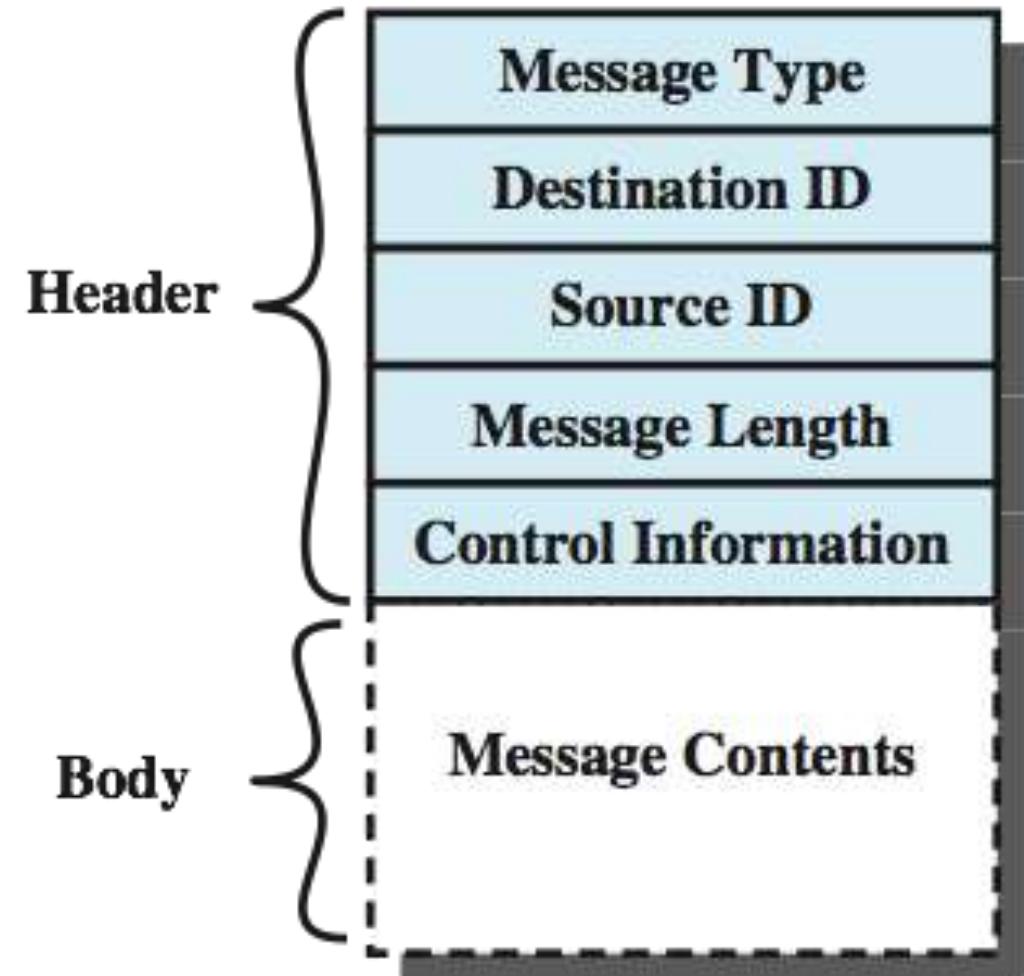
MESSAGE CONSISTS OF TWO PARTS

Header

- Contains information about the message:
e.g...
 - Who sent it?
 - Who is meant to receive it?
 - What type of message is it?

Body

- A block of bytes containing the message itself.
- As with shared memory and streams, it is up to the programmer to structure the attached data in a sensible way.



CONCURRENCY IMPLICATIONS



CAN IPC BE USED AS A MUTUAL EXCLUSION TOOL?

IPC MECHANISMS PROVIDED BY KERNEL BUFFERING...

SYNCHRONISING PROCESSES WITH IPC

- Stream based IPC (pipes, etc.) as well as sockets, provide enforcement of FIFO behavior between readers and writers.
- Together with message queues, they share the following important property:
 - **Data can never be received before it has been sent!**
- This means they can be used for synchronizing concurrent processes, as tools for mutual exclusion.
 - e.g. one process may block waiting to read a message from another process, indicating that it is safe to proceed.
 - The kernel guards access to the IPC resource to enforce mutual exclusion without the need for the programmer to use semaphores/mutexes.

INDIRECT PROCESS COMMUNICATION

STREAMS, MESSAGE QUEUES, ETC.

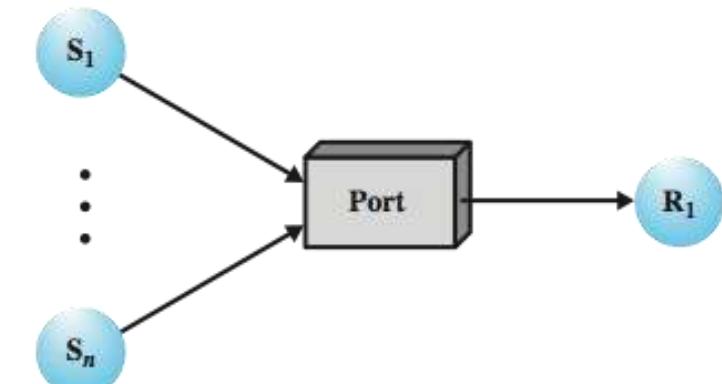
For IPC mechanisms such as pipes, sockets, and message queues, all data must go through the kernel.

This slows things down a little, but enables the kernel to enforce concurrency protections.

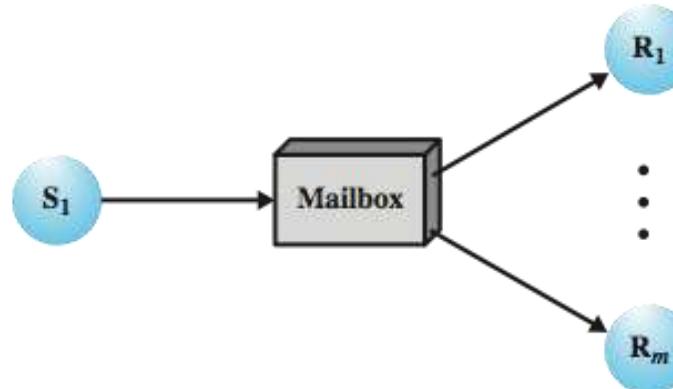
(As discussed earlier.)



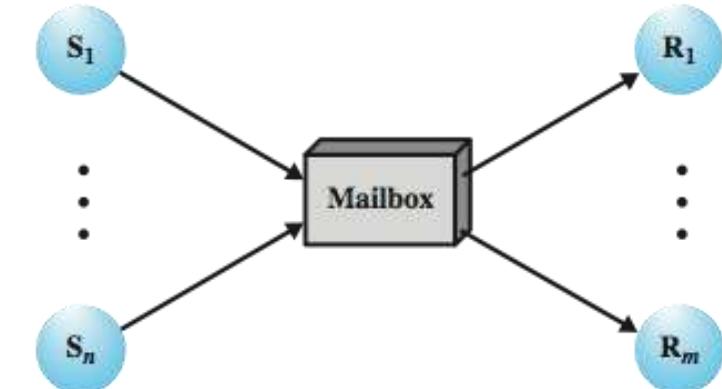
(a) One to one



(b) Many to one



(c) One to many



(d) Many to many

WHAT ABOUT SHARED MEMORY?

NO CONCURRENCY PROTECTIONS AVAILABLE

NOT BUFFERED BY KERNEL

- Shared memory does not pass through buffers in the kernel.
- It is a segment in virtual memory that multiple processes may access directly at any time.
 - Fast, but hazardous!
- Additional synchronization tools (semaphores/mutexes) are needed when dealing with shared memory, to guard access to shared data.

SUMMARY (LECTURE 10)

INTERPROCESS COMMUNICATION

- We have discussed several mechanisms for interprocess communication.
 - Signals are the most primitive
 - Other approaches include pipes, sockets and message queues
- We now have a better understanding of the implications for synchronisation and mutual exclusion, of using IPC tools in an environment where concurrent processes share data.
- **Next week: Security**